

BLG453E - Computer Vision Homework 2

**Segmenting Images, Noise Removal, Perspective Transformation,
and Implementing GANs**

Student: Emil Huseynov

Student ID: 150210906

Course: BLG453E

Assignment: Homework 2

Date: November 18, 2024

Due Date: December 5, 2024

Contents

1	Q1 Segment Image with Otsu's Algorithm (30 pts)	2
1.1	Introduction	2
1.2	Required Functions	2
1.2.1	Function to Add Gaussian Noise	2
1.2.2	Function for Otsu's Thresholding	2
1.3	Q1.1 (10 pts): Gaussian Noise	3
1.3.1	Code	3
1.3.2	Explanation	3
1.3.3	Results	3
1.4	Q1.2 (20 pts): Otsu's Algorithm	4
1.4.1	Code	4
1.4.2	Explanation	4
1.4.3	Results	4
1.5	Analysis	5
2	Q2 Image Noise Addition and Removal (20 pts)	5
2.1	Introduction	5
2.2	Required Functions	5
2.2.1	Function to Add Salt-and-Pepper Noise	5
2.2.2	Function to Add Impulse Noise	5
2.2.3	Function for Median Filtering	5
2.3	Q2.1: Salt-and-Pepper Noise	6
2.3.1	Code	6
2.3.2	Explanation	6
2.3.3	Results	6
2.4	Q2.2: Impulse Noise	6
2.4.1	Code	6
2.4.2	Explanation	7
2.4.3	Results	7
2.5	Analysis	7
3	Q3 Reading the Book Cover (35 pts)	8
3.1	Introduction	8
3.2	Code	8
3.3	Explanation	9
3.4	Results	10
3.5	Analysis	10
4	Understanding and Implementing Generative Adversarial Networks (GANs) (25 pts)	11
4.1	Introduction	11
4.2	Implementation	11
4.2.1	Hyperparameters and Data Loading	11
4.2.2	Generator and Discriminator	11
4.2.3	Initializing Models and Defining Loss Function and Optimizers	12
4.2.4	Training Loop	12
4.2.5	Plotting Results	13
4.3	Analysis of the Generated Images	14
4.3.1	Quality of Generated Images	14
4.3.2	Evolution Over Time	14
4.4	Conclusion	14
4.5	References	14

1 Q1 Segment Image with Otsu's Algorithm (30 pts)

1.1 Introduction

In this part, we aim to segment an image using Otsu's thresholding method. We start by adding Gaussian noise to a simple synthetic image and then apply Otsu's algorithm to find the optimal threshold for segmentation.

1.2 Required Functions

1.2.1 Function to Add Gaussian Noise

Listing 1: Function to Add Gaussian Noise

```
def add_gaussian_noise(image, mean=0, sigma=0.1, clip=True, clip_range=(0, 255)):
    """
    Adds Gaussian noise to an image.
    """
    noise = np.random.normal(mean, sigma, image.shape)
    noisy_image = image + noise
    if clip:
        noisy_image = np.clip(noisy_image, clip_range[0], clip_range[1])
    return noisy_image
```

1.2.2 Function for Otsu's Thresholding

Listing 2: Function for Otsu's Thresholding

```
def otsu_threshold(image):
    # Flatten the image to a 1D array
    pixels = image.flatten()
    total_pixels = pixels.size

    # Compute histogram
    histogram = np.zeros(256)
    unique, counts = np.unique(pixels, return_counts=True)
    for u, c in zip(unique, counts):
        histogram[u] = c

    # Probability of each intensity level
    pdf = histogram / total_pixels

    # Cumulative sum of the probability and mean
    cdf = np.cumsum(pdf)
    cumulative_mean = np.cumsum(pdf * np.arange(256))

    # Total mean intensity of the image
    global_mean = cumulative_mean[-1]

    # Initialize variables
    sigma_b_squared = np.zeros(256)

    # Compute between-class variance for all thresholds
    for T in range(256):
        w0 = cdf[T]
        w1 = 1 - w0

        if w0 == 0 or w1 == 0:
            sigma_b_squared[T] = 0
            continue

        mu0 = cumulative_mean[T] / w0
        mu1 = (global_mean - cumulative_mean[T]) / w1
```

```

sigma_b_squared[T] = w0 * w1 * (mu0 - mu1) ** 2

# Find the optimal threshold
optimal_threshold = np.argmax(sigma_b_squared)

# Segment the image using the optimal threshold
segmented_image = (image >= optimal_threshold).astype(int)

return optimal_threshold, segmented_image

```

1.3 Q1.1 (10 pts): Gaussian Noise

1.3.1 Code

Listing 3: Adding Gaussian Noise to the Image

```

# Create the original image
x = np.zeros((6, 6), dtype=float)
x[0, 0:2] = 1
x[1, 0] = 1
x[3:, 5] = 2
x[4, 4] = 2

# Add Gaussian noise to the image using the function
sigma_noise = 0.2 # Standard deviation of the noise
x_noisy = add_gaussian_noise(x, mean=0, sigma=sigma_noise, clip=True, clip_range=(0, 2))

# Display the original and noisy images
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.imshow(x, cmap='gray', vmin=0, vmax=2)
plt.title('Original_Image')
plt.axis('off')

plt.subplot(1, 2, 2)
plt.imshow(x_noisy, cmap='gray', vmin=0, vmax=2)
plt.title('Image_with_Gaussian_Noise')
plt.axis('off')

```

1.3.2 Explanation

We used a simple 6×6 image with specific pixel values set to 1 and 2, representing different intensity regions. Gaussian noise with a standard deviation of 0.2 was added to the image to simulate real-world noise.

1.3.3 Results




Figure 1: Original Image and Image with Gaussian Noise

1.4 Q1.2 (20 pts): Otsu's Algorithm

1.4.1 Code

Listing 4: Applying Otsu's Algorithm

```
# Prepare the noisy image for Otsu's algorithm
# Scale the noisy image to 8-bit grayscale values (0-255)
x_noisy_scaled = (x_noisy / x_noisy.max()) * 255
x_noisy_scaled = x_noisy_scaled.astype(np.uint8)

# Apply Otsu's thresholding
optimal_threshold, segmented_image = otsu_threshold(x_noisy_scaled)

# Display the segmented image
plt.figure(figsize=(12, 4))

plt.subplot(1, 3, 1)
plt.imshow(x, cmap='gray', vmin=0, vmax=2)
plt.title('Original_Image')
plt.axis('off')

plt.subplot(1, 3, 2)
plt.imshow(x_noisy, cmap='gray', vmin=0, vmax=2)
plt.title('Image_with_Gaussian_Noise')
plt.axis('off')

plt.subplot(1, 3, 3)
plt.imshow(segmented_image, cmap='gray')
plt.title(f'Segmented_Image\nOptimal_T={optimal_threshold}')
plt.axis('off')

plt.tight_layout()
plt.show()

print(f"Optimal_Threshold_(T*):_{optimal_threshold}")
```

1.4.2 Explanation

We scaled the noisy image to 8-bit grayscale values to apply Otsu's algorithm effectively. The optimal threshold was calculated, and the image was segmented based on this threshold.

1.4.3 Results



Figure 2: Original Image, Noisy Image, and Segmented Image

Optimal Threshold (T^*): **103**

1.5 Analysis

The Otsu's algorithm successfully computed an optimal threshold to segment the noisy image into distinct regions. Despite the added Gaussian noise, the segmentation was effective in separating the intensity levels.

2 Q2 Image Noise Addition and Removal (20 pts)

2.1 Introduction

In this part, we added salt-and-pepper noise and impulse noise to the *Lenna.png* image separately. We then applied median filtering to remove the noise from the images.

2.2 Required Functions

2.2.1 Function to Add Salt-and-Pepper Noise

Listing 5: Function to Add Salt-and-Pepper Noise

```
def add_salt_and_pepper_noise(image, salt_prob=0.02, pepper_prob=0.02):
    """
    Adds salt and pepper noise to an image.
    """
    noisy_image = image.copy()
    # Add salt noise
    num_salt = int(salt_prob * image.size)
    salt_coords = [np.random.randint(0, i - 1, num_salt) for i in image.shape[:2]]
    noisy_image[salt_coords[0], salt_coords[1]] = 255

    # Add pepper noise
    num_pepper = int(pepper_prob * image.size)
    pepper_coords = [np.random.randint(0, i - 1, num_pepper) for i in image.shape[:2]]
    noisy_image[pepper_coords[0], pepper_coords[1]] = 0

    return noisy_image
```

2.2.2 Function to Add Impulse Noise

Listing 6: Function to Add Impulse Noise

```
def add_impulse_noise(image, noise_prob=0.1):
    """
    Adds impulse noise (random value noise) to an image.
    """
    noisy_image = image.copy()
    num_noisy_pixels = int(noise_prob * image.size)
    coords = [np.random.randint(0, i - 1, num_noisy_pixels) for i in image.shape[:2]]
    noisy_image[coords[0], coords[1]] = np.random.randint(0, 256, num_noisy_pixels)
    return noisy_image
```

2.2.3 Function for Median Filtering

Listing 7: Function for Median Filtering

```
def median_filter(image, kernel_size=3):
    """
    Applies a median filter to remove salt-and-pepper noise.
    """
    return cv2.medianBlur(image, kernel_size)
```

2.3 Q2.1: Salt-and-Pepper Noise

2.3.1 Code

Listing 8: Adding Salt-and-Pepper Noise and Median Filtering

```
# Read the PNG image
image = cv2.imread("Lenna.png", cv2.IMREAD_GRAYSCALE)

# Add salt-and-pepper noise
noisy_image = add_salt_and_pepper_noise(image, 0.03, 0.03)

# Denoise using median filtering
denoised_image = median_filter(noisy_image)

# Display the images
plt.figure(figsize=(12, 6))
plt.subplot(1, 3, 1), plt.imshow(image, cmap='gray'), plt.title("Original_Image"), plt.axis('off')
plt.subplot(1, 3, 2), plt.imshow(noisy_image, cmap='gray'), plt.title("Noisy_Image_(Salt_and_Pepper)"), plt.axis('off')
plt.subplot(1, 3, 3), plt.imshow(denoised_image, cmap='gray'), plt.title("Denoised_Image"), plt.axis('off')
plt.show()
```

2.3.2 Explanation

We introduced salt-and-pepper noise with a probability of 0.03 for both salt and pepper. Median filtering with a kernel size of 3 was applied to remove the noise.

2.3.3 Results




Figure 3: Original Image, Noisy Image (Salt-and-Pepper), and Denoised Image

2.4 Q2.2: Impulse Noise

2.4.1 Code

Listing 9: Adding Impulse Noise and Median Filtering

```
# Read the PNG image
image = cv2.imread("Lenna.png", cv2.IMREAD_GRAYSCALE)

# Add impulse noise
noisy_image = add_impulse_noise(image, 0.15)
```

```

# Denoise using median filtering
denoised_image = median_filter(noisy_image)

# Display the images
plt.figure(figsize=(12, 6))
plt.subplot(1, 3, 1), plt.imshow(image, cmap='gray'), plt.title("Original_Image"), plt.axis('off')
plt.subplot(1, 3, 2), plt.imshow(noisy_image, cmap='gray'), plt.title("Noisy_Image_(Impulse_Noise)"), plt.axis('off')
plt.subplot(1, 3, 3), plt.imshow(denoised_image, cmap='gray'), plt.title("Denoised_Image"), plt.axis('off')
plt.show()

```

2.4.2 Explanation

Impulse noise was added with a probability of 0.15. Median filtering effectively reduced the noise, enhancing image quality.

2.4.3 Results




Figure 4: Original Image, Noisy Image (Impulse Noise), and Denoised Image

2.5 Analysis

Median filtering proved to be effective in removing both salt-and-pepper noise and impulse noise. The filter replaces each pixel value with the median of neighboring pixel values, preserving edges while reducing noise.

3 Q3 Reading the Book Cover (35 pts)

3.1 Introduction

In this part, we aimed to detect the corners of a book in an image and apply a perspective transformation to obtain a front-facing view of the book cover. This process involves edge detection, contour detection, and perspective warping.

3.2 Code

Listing 10: Perspective Transformation to Read Book Cover

```
# Step 1: Load the Image
image_path = "gogol.jpg"
image = cv2.imread(image_path)
image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Convert to RGB for visualization

# Step 2: Preprocess the Image
# Convert to grayscale
gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

# Apply GaussianBlur to reduce noise
blurred = cv2.GaussianBlur(gray, (5, 5), 0)

# Use Canny edge detection
edges = cv2.Canny(blurred, threshold1=50, threshold2=150)

# Step 3: Detect Contours
contours, _ = cv2.findContours(edges, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)

# Filter for the largest rectangular contour
book_contour = None
for contour in contours:
    epsilon = 0.02 * cv2.arcLength(contour, True)
    approx = cv2.approxPolyDP(contour, epsilon, True)
    if len(approx) == 4: # Look for quadrilaterals
        book_contour = approx
        break

# Check if a contour was found and draw it on the image
if book_contour is not None:
    detected_image = image_rgb.copy()
    cv2.drawContours(detected_image, [book_contour], -1, (255, 0, 0), 3)

# Step 4: Order the corner points
points = book_contour.reshape(4, 2)
rect = np.zeros((4, 2), dtype="float32")
s = points.sum(axis=1)
rect[0] = points[np.argmin(s)] # Top-left
rect[2] = points[np.argmax(s)] # Bottom-right
diff = np.diff(points, axis=1)
rect[1] = points[np.argmin(diff)] # Top-right
rect[3] = points[np.argmax(diff)] # Bottom-left

print("Ordered_corner_points:", rect)

# Step 5: Define destination points
width = 400 # Desired width
height = 600 # Desired height
destination_points = np.array([
    [0, 0],
    [width - 1, 0],
    [width - 1, height - 1],
    [0, height - 1]
], dtype="float32")
```

```

# Step 6: Compute perspective transformation matrix
matrix = cv2.getPerspectiveTransform(rect, destination_points)

# Apply perspective warp
warped_image = cv2.warpPerspective(image_rgb, matrix, (width, height))

# Plot all intermediate and final results together
fig, axes = plt.subplots(1, 3, figsize=(10, 6))
axes = axes.ravel()

# Original image
axes[0].imshow(image_rgb)
axes[0].set_title("Original_Image")
axes[0].axis("off")

# Detected contour
if book_contour is not None:
    axes[1].imshow(detected_image)
    axes[1].set_title("Detected_Contour")
else:
    axes[1].imshow(image_rgb)
    axes[1].set_title("No_Contour_Detected")
axes[1].axis("off")

# Perspective corrected image
if book_contour is not None:
    axes[2].imshow(warped_image)
    axes[2].set_title("Perspective_Corrected_Image")
else:
    axes[2].imshow(image_rgb)
    axes[2].set_title("No_Perspective_Correction")
axes[2].axis("off")

plt.tight_layout()
plt.show()

```

3.3 Explanation

1. **Preprocessing:** Converted the image to grayscale and applied Gaussian blur to reduce noise.
2. **Edge Detection:** Used the Canny edge detector to find edges in the image.
3. **Contour Detection:** Found contours in the edged image and approximated them to polygons.
4. **Corner Detection:** Identified contours that are quadrilaterals, assuming they correspond to the book.
5. **Perspective Transformation:** Computed a transformation matrix to warp the image to a top-down view of the book cover.

3.4 Results




Figure 5: Original Image, Detected Contour, and Perspective Corrected Image

3.5 Analysis

The process successfully detected the book's corners and corrected the perspective, allowing for a clear, front-facing view of the book cover. This technique is useful in document scanning and object recognition applications.

4 Understanding and Implementing Generative Adversarial Networks (GANs) (25 pts)

4.1 Introduction

Generative Adversarial Networks (GANs) are powerful models capable of generating realistic data. In this task, we implemented a simple GAN using PyTorch and trained it on the CIFAR-10 dataset.

4.2 Implementation

4.2.1 Hyperparameters and Data Loading

Listing 11: Hyperparameters and Data Loading

```
# Hyperparameters
batch_size = 128
image_size = 32
channels_img = 3 # CIFAR-10 images are RGB
z_dim = 100 # Size of the noise vector (input to Generator)
num_epochs = 25
learning_rate = 0.0002

# Device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Dataset and DataLoader
transform = transforms.Compose([
    transforms.Resize(image_size),
    transforms.ToTensor(),
    transforms.Normalize([0.5]*channels_img, [0.5]*channels_img) # Normalize images to [-1, 1]
])

train_dataset = torchvision.datasets.CIFAR10(
    root='./data', train=True, download=True, transform=transform
)

dataloader = torch.utils.data.DataLoader(
    train_dataset, batch_size=batch_size, shuffle=True
)
```

4.2.2 Generator and Discriminator

Listing 12: Generator and Discriminator Networks

```
# Generator
class Generator(nn.Module):
    def __init__(self, z_dim, channels_img):
        super(Generator, self).__init__()
        self.main = nn.Sequential(
            nn.ConvTranspose2d(z_dim, 512, 4, 1, 0, bias=False), # Output: (512, 4, 4)
            nn.BatchNorm2d(512),
            nn.ReLU(True),
            nn.ConvTranspose2d(512, 256, 4, 2, 1, bias=False), # Output: (256, 8, 8)
            nn.BatchNorm2d(256),
            nn.ReLU(True),
            nn.ConvTranspose2d(256, 128, 4, 2, 1, bias=False), # Output: (128, 16, 16)
            nn.BatchNorm2d(128),
            nn.ReLU(True),
            nn.ConvTranspose2d(128, channels_img, 4, 2, 1, bias=False), # Output: (3, 32, 32)
            nn.Tanh()
        )

    def forward(self, input):
        return self.main(input)
```

```

# Discriminator
class Discriminator(nn.Module):
    def __init__(self, channels_img):
        super(Discriminator, self).__init__()
        self.main = nn.Sequential(
            nn.Conv2d(channels_img, 128, 4, 2, 1, bias=False),      # Output: (128, 16, 16)
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(128, 256, 4, 2, 1, bias=False),              # Output: (256, 8, 8)
            nn.BatchNorm2d(256),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(256, 512, 4, 2, 1, bias=False),              # Output: (512, 4, 4)
            nn.BatchNorm2d(512),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(512, 1, 4, 1, 0, bias=False),                # Output: (1, 1, 1)
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.main(input).view(-1, 1).squeeze(1)

```

4.2.3 Initializing Models and Defining Loss Function and Optimizers

Listing 13: Initialization and Optimization Setup

```

# Initialize models
netG = Generator(z_dim, channels_img).to(device)
netD = Discriminator(channels_img).to(device)

# Weight initialization
def weights_init(m):
    classname = m.__class__.__name__
    if classname.find('Conv') != -1 or classname.find('BatchNorm') != -1:
        nn.init.normal_(m.weight.data, 0.0, 0.02)

netG.apply(weights_init)
netD.apply(weights_init)

# Loss function and optimizers
criterion = nn.BCELoss() # Binary Cross Entropy Loss

fixed_noise = torch.randn(64, z_dim, 1, 1, device=device) # For generating consistent images

optimizerD = optim.Adam(netD.parameters(), lr=learning_rate, betas=(0.5, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=learning_rate, betas=(0.5, 0.999))

```

4.2.4 Training Loop

Listing 14: Training Loop for GAN

```

# Training Loop
img_list = []
G_losses = []
D_losses = []

print("Starting Training Loop...")
for epoch in range(num_epochs):
    for i, data in enumerate(dataloader, 0):
        # (1) Update D network
        netD.zero_grad()
        real_images = data[0].to(device)
        b_size = real_images.size(0)
        label = torch.full((b_size,), 1., dtype=torch.float, device=device)

        output = netD(real_images)

```

```

errD_real = criterion(output, label)
errD_real.backward()
D_x = output.mean().item()

# Generate fake images
noise = torch.randn(b_size, z_dim, 1, 1, device=device)
fake_images = netG(noise)
label.fill_(0.)

output = netD(fake_images.detach())
errD_fake = criterion(output, label)
errD_fake.backward()
D_G_z1 = output.mean().item()

errD = errD_real + errD_fake
optimizerD.step()

# (2) Update G network
netG.zero_grad()
label.fill_(1.)

output = netD(fake_images)
errG = criterion(output, label)
errG.backward()
D_G_z2 = output.mean().item()

optimizerG.step()

# Output training stats
if i % 100 == 0:
    print(
        f'[{epoch}/{num_epochs}] [{i}/{len(dataloader)}]\t'
        f'Loss_D: {errD.item():.4f}\tLoss_G: {errG.item():.4f}\t'
        f'D(x): {D_x:.4f}\tD(G(z)): {D_G_z1:.4f}/{D_G_z2:.4f}'
    )

# Save Losses for plotting later
G_losses.append(errG.item())
D_losses.append(errD.item())

# Save generated images for visualization
with torch.no_grad():
    fake = netG(fixed_noise).detach().cpu()
    img_list.append(torchvision.utils.make_grid(fake, padding=2, normalize=True))

```

4.2.5 Plotting Results

Listing 15: Plotting Real and Fake Images

```

# Function to show images
def show_images(images, title):
    plt.figure(figsize=(48,16))
    plt.axis("off")
    plt.title(title)
    plt.imshow(np.transpose(images, (1,2,0)))
    plt.show()

# Plot the real images
real_batch = next(iter(dataloader))
real_images = torchvision.utils.make_grid(real_batch[0][:64], padding=2, normalize=True)
show_images(real_images.cpu(), "Real_Images")

# Plot the fake images from the last epoch
fake_images = img_list[-1]
show_images(fake_images, "Fake_Images")

```

4.3 Analysis of the Generated Images

4.3.1 Quality of Generated Images

Before Training: The Generator produces images resembling random noise.

Midway Through Training: Images start to show basic structures and colors similar to CIFAR-10 images but lack fine details.

After Full Training: Generated images are more detailed and closely resemble real images from the dataset.

4.3.2 Evolution Over Time

The Generator learns to produce images that increasingly fool the Discriminator, improving the realism of the generated images. The adversarial training leads both networks to enhance their performance iteratively.

4.4 Conclusion

By implementing and training this GAN, we observed how the Generator and Discriminator networks improve through competition. The Generator learns to produce more realistic images, while the Discriminator becomes better at distinguishing real from fake, showcasing the effectiveness of adversarial training in generating synthetic data.

4.5 References

References

- [1] Goodfellow, Ian, et al. "Generative adversarial nets." *Advances in neural information processing systems* 27 (2014).